

Intensity Transformations and Neighborhood Filtering

EN3160 - Image Processing and Machine Vision

Balasooriya B.A.P.I.
220054N

© PankajaBalasooriya/EN3160_Image_Processing_and_Machine_Vision

 View Jupyter Notebook

Question 1: Basic Intensity Transformation

Piecewise intensity transformation function is constructed by linearly interpolating between the control points "c".

```

1 c = np.array([(50, 50), (50, 100), (150, 255), (150,150)])
2
3 t1 = np.linspace(0, c[0,1], c[0,0] + 1 -0).astype('uint8')
4 t2 = np.linspace(c[0,1]+1, c[1,1], c[1,0] - c[0,0]).astype('uint8')
5 t3 = np.linspace(c[1,1]+1, c[2,1], c[2,0] - c[1,0]).astype('uint8')
6 t4 = np.linspace(c[2,1]+1, c[3,1], c[3,0] - c[2,0]).astype('uint8')
7 t5 = np.linspace(c[3,1]+1, 255, 255 - c[3,0]).astype('uint8')
8
9 transform = np.concatenate((t1, t2), axis=0).astype('uint8')
10 transform = np.concatenate((transform, t3), axis=0).astype('uint8')
11 transform = np.concatenate((transform, t4), axis=0).astype('uint8')
12 transform = np.concatenate((transform, t5), axis=0).astype
13
14 img_emma_transformed = cv.LUT(img_emma, transform)

```



(a) Original Image



(b) Transformed Image

The transformation creates jump discontinuities that result in high contrast regions. Mid-intensity pixels are enhanced while preserving the extreme intensity values. This leads to improved visibility of features in the mid-range while maintaining the overall structure of the image.

Figure 1: Intensity transformation results

Question 2: Brain Image Enhancement

```

1 # White Matter Enhancement transformation
2 c = np.array([(140, 0), (180, 180)])
3 t1 = np.zeros(c[0, 0] + 1).astype('uint8')
4 t2 = np.linspace(c[0, 0] + 1, c[1, 1], c[1, 0] - c[0, 0]).astype('uint8')
5 t3 = np.zeros(256 - len(t1) - len(t2)).astype('uint8')
6 intensity_transform_white_matter_lut = np.concatenate((t1, t2, t3), axis=0).astype('uint8')
7 img_brain_white_matter = cv.LUT(img_brain, intensity_transform_white_matter_lut)

```

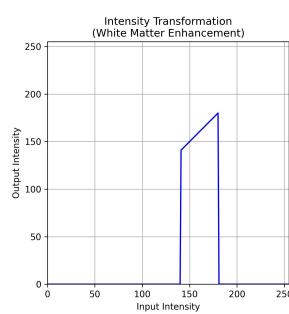
```

1 # Grey Matter Enhancement transformation
2 c = np.array([(180, 0), (210, 210)])
3 t1 = np.zeros(c[0, 0] + 1).astype('uint8')
4 t2 = np.linspace(c[0, 0] + 1, c[1, 1], c[1, 0] - c[0, 0]).astype('uint8')
5 t3 = np.zeros(256 - len(t1) - len(t2)).astype('uint8')
6 intensity_transform_grey_matter_lut = np.concatenate((t1, t2, t3)).astype('uint8')
7 img_brain_grey_matter = cv.LUT(img_brain, intensity_transform_grey_matter_lut)

```



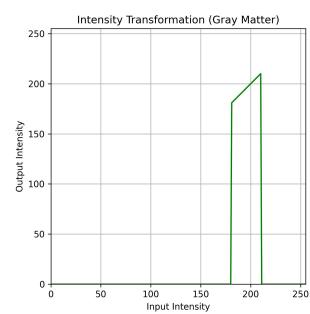
(a) White Matter Enhanced



(b) White Matter Plot



(c) Gray Matter Enhanced



(d) Gray Matter Plot

Figure 2: White and Gray matter enhancement results and transformation curves

The intensity transformations offer smooth contrast enhancement within specific intensity ranges, based on the color values of white and gray matter, effectively enhancing the visibility of these tissues.

Question 3: Gamma Correction

L*a*b* Color Space Conversion

```

1 img_lab = cv.cvtColor(img_bgr, cv.COLOR_BGR2Lab)
2 L, a, b = cv.split(img_lab)

```

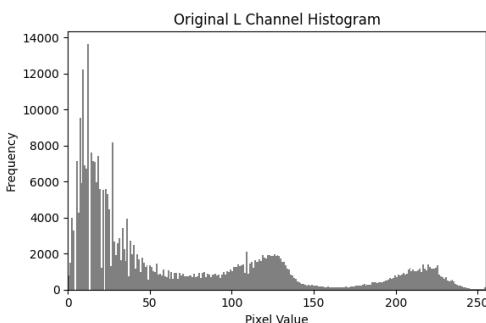
Applying γ value correction

```

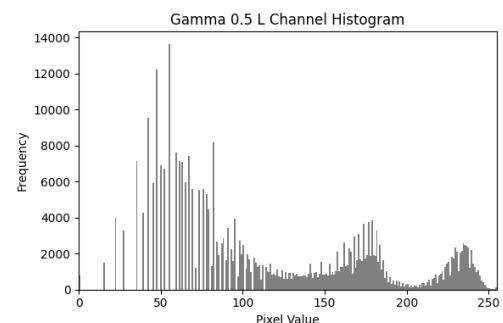
1 L_norm = L.astype(np.float32) / 255.0
2 # Apply gamma correction
3 gamma = 0.5 # <1 brightens, >1 darkens
4 L_gamma = np.power(L_norm, gamma)
5 L_new = np.clip(L_gamma * 255.0, 0, 255).astype(np.uint8)
6 # Merge and convert back to BGR
7 img_lab_new = cv.merge((L_new, a, b))
8 img_bgr_new = cv.cvtColor(img_lab_new, cv.COLOR_Lab2BGR)

```

γ value is chosen as 0.5 by comparing the images by changing the values of γ , which lightens the shadows while preserving highlights, improving overall image visibility.



(a) Original L plane Histogram



(b) Gamma Corrected L plane Histogram

Question 4: Vibrance Enhancement

HSV Decomposition

```
1 spider_hsv = cv.cvtColor(spider_img, cv.COLOR_BGR2HSV)
2 h, s, v = cv.split(spider_hsv)
```



(a) Hue



(b) Saturation



(c) Value

Figure 4: HSV plane decomposition

Vibrance Transformation

```
1 def transformation(x: np.ndarray, a: float, sigma: int=70) -> np.ndarray:
2     x_float = x.astype(np.float64) # To stop overflow
3     exponent = -1 * (x_float - 128)**2 / (2 * sigma**2)
4     fx = x_float + a * 128 * np.exp(exponent)
5     return np.clip(fx, 0, 255).astype(np.uint8)
```

Applying the transformation

```
1 a_val = 0.7
2 sigma_val = 70
3 s_transformed = transformation(s, a=a_val, sigma=sigma_val)
4 spider_hsv_transformed = cv.merge([h, s_transformed, v])
5 spider_transformed_bgr = cv.cvtColor(spider_hsv_transformed, cv.COLOR_HSV2BGR)
```

A value of $a = 0.7$ was chosen to provide a balanced enhancement of image vibrance without introducing oversaturation or unnatural colors. At this level, colorful regions are noticeably intensified, improving visual appeal, while grayscale and low-saturation regions remain largely unaffected. This ensures that the transformation selectively enhances vivid areas without distorting the overall color balance, making it suitable for natural-looking vibrance enhancement.

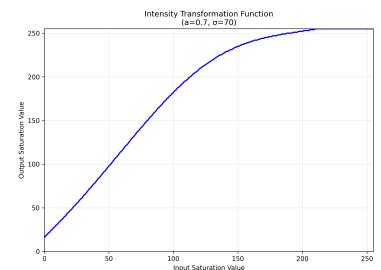
Results Comparison



(a) Original Image



(b) Vibrance-enhanced image



(c) Intensity Transformation

Figure 5: HSV plane decomposition

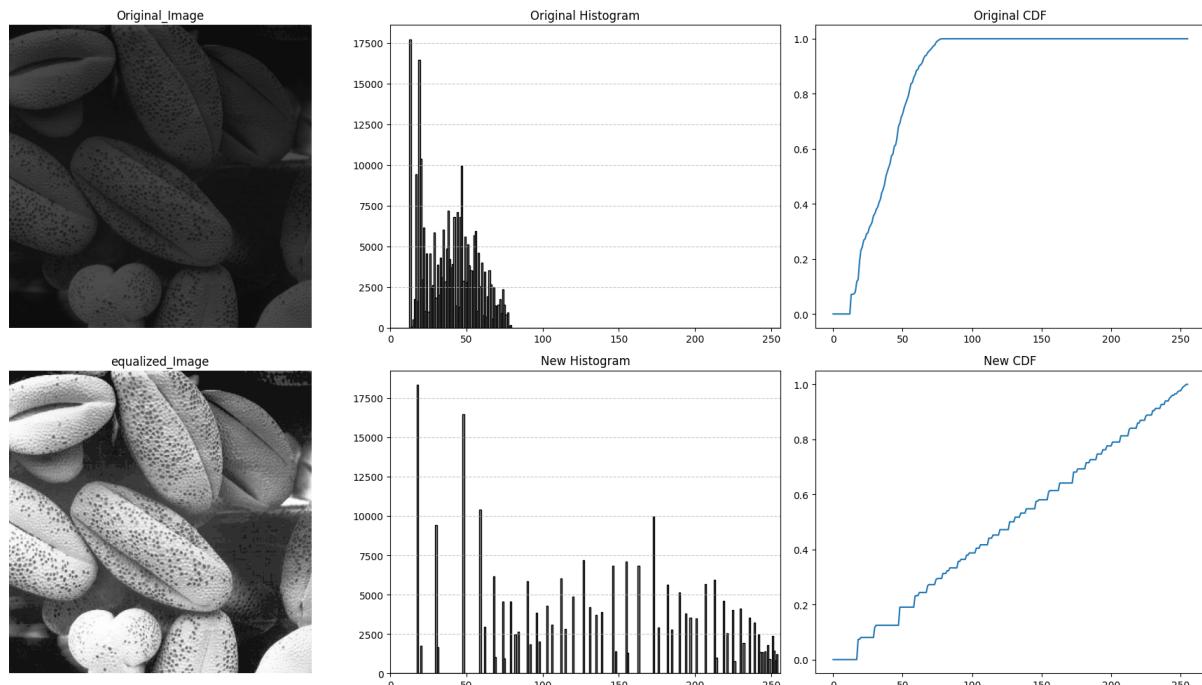
Question 5: Histogram Equalization

```

1 def calculate_histogram(image: np.ndarray, height: int, width) -> np.ndarray:
2     # Calculate histogram
3     hist = np.zeros(256)
4     for i in range(height):
5         for j in range(width):
6             hist[image[i, j]] += 1
7     return hist
8
9 def calculate_cdf(image: np.ndarray, height, width):
10    # Calculate cumulative distribution function (CDF)
11    total_pixels = height * width
12    hist = calculate_histogram(image, height, width)
13    cdf = np.zeros(256)
14    cdf[0] = hist[0]
15    for i in range(1, 256):
16        cdf[i] = cdf[i-1] + hist[i]
17    # Normalize CDF
18    cdf_normalized = cdf / total_pixels
19    return cdf_normalized, hist
20
21 def histogram_equalization(image: np.ndarray) -> np.ndarray:
22     height, width = image.shape
23     total_pixels = height * width
24     original_cdf, original_histogram = calculate_cdf(image, height, width)
25     # Apply transformation
26     equalized_image = np.zeros_like(image)
27     for i in range(height):
28         for j in range(width):
29             equalized_image[i, j] = (255 * original_cdf[image[i, j]]).astype(np.uint8)
30     return equalized_image
31
32
33 img = cv.imread("./alimages/shells.tif", cv.IMREAD_GRAYSCALE)
34 equalized_img = histogram_equalization(img, debug=True)

```

Results and Analysis



Question 6: Foreground Histogram Equalization

HSV Analysis and Masking

The Saturation channel is selected as it provides clear separation of the foreground based on color intensity, effectively isolating the subject.



(a) Hue



(b) Saturation



(c) Value

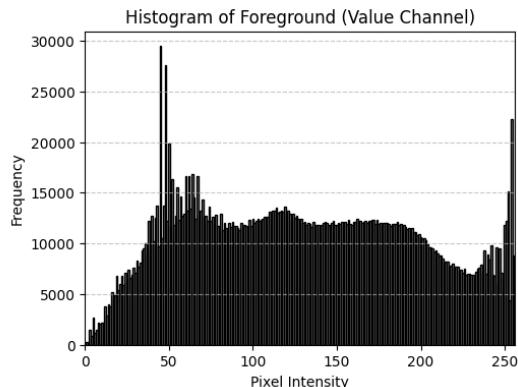
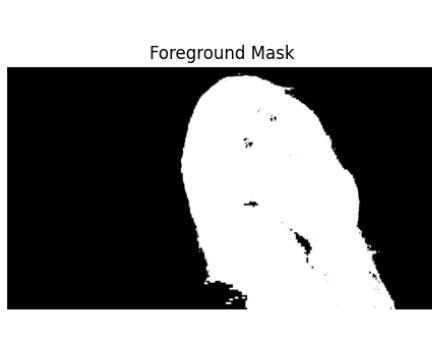
Figure 6: HSV plane decomposition

Foreground Extraction

```

1 threshold_value = 13
2 _, mask = cv.threshold(s, threshold_value, 255, cv.THRESH_BINARY)
3
4 foreground = cv.bitwise_and(jeniffer_img, jeniffer_img, mask=mask)
5 # Compute histogram of the value channel for foreground pixels only
6 foreground_v = cv.bitwise_and(v, v, mask=mask)

```



Selective Histogram Equalization

Application of histogram equalization only to the foreground.

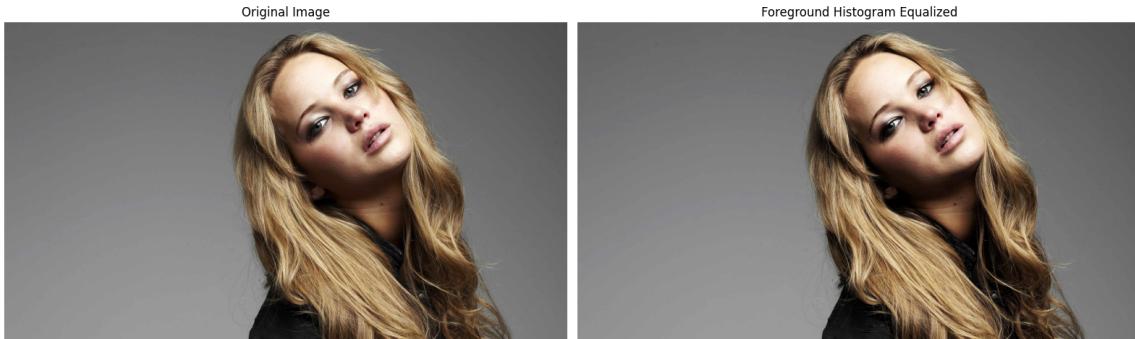
```

1 cumsum = np.cumsum(hist_original)
2 cumsum_normalized = cumsum * 255 / cumsum[-1] # Normalize cumulative sum
3
4 lut = np.zeros(256, dtype=np.uint8) # Create lookup table
5 for i in range(256):
6     lut[i] = int(cumsum_normalized[i])
7
8 # Apply histogram equalization only to foreground pixels
9 v_equalized = v.copy()
10 v_equalized[mask > 0] = lut[v[mask > 0]]
11
12 hsv_equalized = cv.merge([h, s, v_equalized])

```

Results

Selective equalization enhances foreground details while preserving the background, resulting in improved contrast and clearer distinction of the subject.



Question 7: Sobel Filtering

Method 1: Using filter2D

```

1 sobel_x = np.array([[-1, -2, -1],
2                     [0, 0, 0],
3                     [1, 2, 1]])
4
5 sobel_y = np.array([[-1, 0, 1],
6                     [-2, 0, 2],
7                     [-1, 0, 1]])
8
9 einstein_sobel_x = cv.filter2D(einstein_img, cv.CV_64F, sobel_x)
0 einstein_sobel_y = cv.filter2D(einstein_img, cv.CV_64F, sobel_y)

```

Method 2: Custom Implementation

```

1 def sobel_convolution(image, kernel):
2     h, w = image.shape
3     kh, kw = kernel.shape
4     pad_h, pad_w = kh // 2, kw // 2
5     padded = np.pad(image, ((pad_h, pad_h), (pad_w, pad_w)), mode='constant')
6     output = np.zeros_like(image, dtype=np.float64)
7     for i in range(h):
8         for j in range(w):
9             region = padded[i:i+kh, j:j+kw]
10            output[i, j] = np.sum(region * kernel)
11
12    return output
13
14 sobel_x_manual = sobel_convolution(einstein_img, sobel_x)
15 sobel_y_manual = sobel_convolution(einstein_img, sobel_y)

```

Method 3: Separable Filters

```

1 kx = np.array([1, 0, -1], dtype=np.float32).reshape(1, 3)
2 ky = np.array([1, 2, 1], dtype=np.float32).reshape(3, 1)
3 # First filter horizontally, then vertically
4 temp = cv.filter2D(einstein_img, cv.CV_64F, kx)
5 sobel_x_sep = cv.filter2D(temp, cv.CV_64F, ky)

```

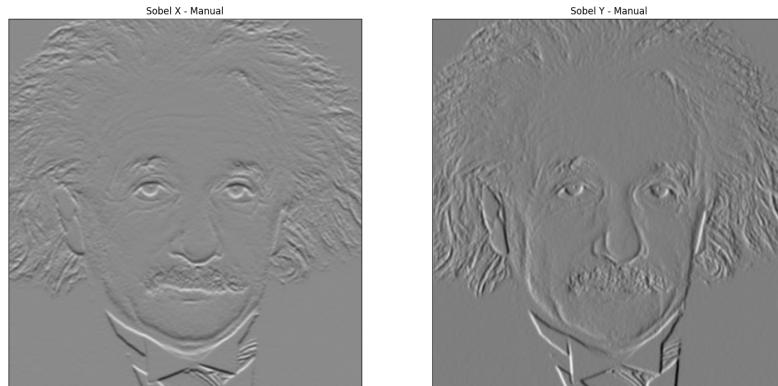


Figure 7: Sobel Filtering

Question 8: Image Zooming

0.0.1 Interpolation Methods

```

1 def zoom_image(img, target_size=None, scale=None, method="nearest"):
2     if method == "nearest":
3         interp = cv.INTER_NEAREST
4     elif method == "bilinear":
5         interp = cv.INTER_LINEAR
6
7     if target_size is not None:
8         return cv.resize(img, target_size, interpolation=interp)
9     elif scale is not None and (0 < scale <= 10):
10        h, w = img.shape[:2]
11        return cv.resize(img, (int(w * scale), int(h * scale)), interpolation=interp)
12
13 def normalized_ssd(img1, img2):
14     diff = img1.astype(np.float32) - img2.astype(np.float32)
15     ssd = np.sum(diff ** 2)
16     norm_ssd = ssd / np.sum(img1.astype(np.float32) ** 2)
17
18     return norm_ssd

```



Results Analysis

Bilinear interpolation produces smoother and more visually pleasing results compared to nearest neighbor. Additionally, the slightly lower SSD value for bilinear indicates a closer match to the reference image.

Question 9: Flower

```

1 mask = np.zeros(img.shape[:2], np.uint8)
2 bkgModel = np.zeros((1, 65), np.float64)
3 fgdModel = np.zeros((1, 65), np.float64)
4 rect = (50, 50, img.shape[1] - 50, img.shape[0] - 50)
5 # Run GrabCut
6 cv.grabCut(img_rgb, mask, rect, bkgModel, 5, cv.GC_INIT_WITH_RECT)
7 mask2 = np.where((mask == cv.GC_FGD) | (mask == cv.GC_PR_FGD), 1, 0).astype('uint8')
8 foreground = img_rgb * mask2[:, :, np.newaxis]
9 background = img_rgb * (1 - mask2[:, :, np.newaxis])
0 segmentation_mask = (mask2 * 255).astype(np.uint8)
1 # Blur the entire image
2 blurred = cv.GaussianBlur(img_rgb, (25, 25), 0)
3 enhanced = blurred.copy()
4 enhanced[mask2 == 1] = img_rgb[mask2 == 1]

```



Reason for dark background at flower edges

When applying Gaussian blur to the background, the kernel averages surrounding pixels. Near the flower edges, some pixels are zeroed out from foreground extraction, causing the averaged values to be darker than the original background.